

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Importing orders table

```
[3]: orders = pd.read_csv('Orders.csv')
```

```
[4]: orders.describe()
```

```
[4]:
```

	OrderID	CustomerID	ProductID	Quantity	TotalAmount	StoreID
count	10000.00000	10000.00000	10000.000000	10000.000000	10000.000000	10000.000000
mean	8000.50000	11000.50000	1024.915200	2.981500	501.416700	2049.074200
std	2886.89568	2886.89568	14.610303	1.413421	431.949384	28.908574
min	3001.00000	6001.00000	1000.000000	1.000000	14.000000	2000.000000
25%	5500.75000	8500.75000	1012.000000	2.000000	180.000000	2024.000000
50%	8000.50000	11000.50000	1025.000000	3.000000	385.000000	2049.000000
75%	10500.25000	13500.25000	1038.000000	4.000000	700.000000	2074.000000
max	13000.00000	16000.00000	1050.000000	5.000000	2995.000000	2099.000000

```
[5]: orders.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   OrderID          10000 non-null   int64  
 1   CustomerID       10000 non-null   int64  
 2   OrderDateTime    10000 non-null   object  
 3   ProductID        10000 non-null   int64  
 4   Quantity         10000 non-null   int64  
 5   TotalAmount      10000 non-null   int64  
 6   StoreID          10000 non-null   int64  
 7   OrderStatus       10000 non-null   object  
 8   ExpectedDeliveryTime 10000 non-null   object  
 9   ActualDeliveryTime 10000 non-null   object  
dtypes: int64(6), object(4)
memory usage: 781.4+ KB
```

```
[6]: orders.isnull().sum()
```

```
[6]:
```

OrderID	0
CustomerID	0
OrderDateTime	0
ProductID	0
Quantity	0
TotalAmount	0
StoreID	0
OrderStatus	0
ExpectedDeliveryTime	0
ActualDeliveryTime	0

```
dtype: int64
```

```
[7]: orders.dtypes
```

```
[7]:
```

OrderID	int64
CustomerID	int64
OrderDateTime	object
ProductID	int64
Quantity	int64
TotalAmount	int64
StoreID	int64
OrderStatus	object
ExpectedDeliveryTime	object
ActualDeliveryTime	object

```
dtype: object
```

```
[8]: orders.duplicated()
```

```
[8]: 0      False
1      False
2      False
3      False
4      False
...
9995    False
9996    False
9997    False
9998    False
9999    False
Length: 10000, dtype: bool
```



Changing datatypes of orderdatetime from object to datetime

```
[9]: orders['OrderDateTime'] = pd.to_datetime(orders['OrderDateTime'], format='%Y-%m-%d %H:%M:%S.%f')
orders
```

	OrderID	CustomerID	OrderDateTime	ProductID	Quantity	TotalAmount	StoreID	OrderStatus	ExpectedDeliveryTime	ActualDeliveryTime
0	3001	6001	2025-01-22 20:02:03.428463	1025	2	90	2051	Delivered	2025-01-22 20:12:03.428463	2025-01-22 20:27:03.428463
1	3002	6002	2025-02-18 15:58:03.428463	1027	1	99	2084	Delivered	2025-02-18 16:08:03.428463	2025-02-18 16:10:03.428463
2	3003	6003	2025-02-04 06:58:03.428463	1001	3	780	2008	Delivered	2025-02-04 07:08:03.428463	2025-02-04 07:10:03.428463
3	3004	6004	2025-02-08 17:00:03.428463	1025	4	180	2008	Delivered	2025-02-08 17:10:03.428463	2025-02-08 17:25:03.428463
4	3005	6005	2025-02-06 18:43:03.428463	1023	3	405	2071	Delivered	2025-02-06 18:53:03.428463	2025-02-06 19:08:03.428463
...
9995	12996	15996	2025-02-06 17:44:03.428463	1035	1	90	2036	Delivered	2025-02-06 17:54:03.428463	2025-02-06 17:59:03.428463
9996	12997	15997	2025-02-08 20:34:03.428463	1043	4	560	2035	Delivered	2025-02-08 20:44:03.428463	2025-02-08 20:54:03.428463
9997	12998	15998	2025-02-08 06:43:03.428463	1017	4	960	2085	Delivered	2025-02-08 06:53:03.428463	2025-02-08 06:55:03.428463
9998	12999	15999	2025-02-15 05:58:03.428463	1039	4	440	2046	Delivered	2025-02-15 06:08:03.428463	2025-02-15 06:08:03.428463
9999	13000	16000	2025-01-21 18:20:03.428463	1009	4	140	2008	Delivered	2025-01-21 18:30:03.428463	2025-01-21 18:30:03.428463

10000 rows × 10 columns

```
[10]: orders.dtypes
```

```
[10]: OrderID          int64
CustomerID        int64
OrderDateTime      datetime64[ns]
ProductID          int64
Quantity           int64
TotalAmount         int64
StoreID            int64
OrderStatus         object
ExpectedDeliveryTime  object
ActualDeliveryTime  object
dtype: object
```



```
[11]: orders['ExpectedDeliveryTime'] = pd.to_datetime(orders['ExpectedDeliveryTime'], format='%Y-%m-%d %H:%M:%S.%f')
orders
orders['ActualDeliveryTime'] = pd.to_datetime(orders['ActualDeliveryTime'], format='%Y-%m-%d %H:%M:%S.%f')
orders
```

	OrderID	CustomerID	OrderDateTime	ProductID	Quantity	TotalAmount	StoreID	OrderStatus	ExpectedDeliveryTime	ActualDeliveryTime
0	3001	6001	2025-01-22 20:02:03.428463	1025	2	90	2051	Delivered	2025-01-22 20:12:03.428463	2025-01-22 20:27:03.428463
1	3002	6002	2025-02-18 15:58:03.428463	1027	1	99	2084	Delivered	2025-02-18 16:08:03.428463	2025-02-18 16:10:03.428463
2	3003	6003	2025-02-04 06:58:03.428463	1001	3	780	2008	Delivered	2025-02-04 07:08:03.428463	2025-02-04 07:10:03.428463

	CustomerID	ProductID	OrderDate	UnitPrice	Quantity	Freight	ShipVia	ShippedOn	DeliveredOn	CreatedOn	UpdatedOn
3	3004	6004	2025-02-08 17:00:03.428463	1025	4	180	2008	Delivered	2025-02-08 17:10:03.428463	2025-02-08 17:25:03.428463	
4	3005	6005	2025-02-06 18:43:03.428463	1023	3	405	2071	Delivered	2025-02-06 18:53:03.428463	2025-02-06 19:08:03.428463	
...
9995	12996	15996	2025-02-06 17:44:03.428463	1035	1	90	2036	Delivered	2025-02-06 17:54:03.428463	2025-02-06 17:59:03.428463	
9996	12997	15997	2025-02-08 20:34:03.428463	1043	4	560	2035	Delivered	2025-02-08 20:44:03.428463	2025-02-08 20:54:03.428463	
9997	12998	15998	2025-02-08 06:43:03.428463	1017	4	960	2085	Delivered	2025-02-08 06:53:03.428463	2025-02-08 06:55:03.428463	
9998	12999	15999	2025-02-15 05:58:03.428463	1039	4	440	2046	Delivered	2025-02-15 06:08:03.428463	2025-02-15 06:08:03.428463	
9999	13000	16000	2025-01-21 18:20:03.428463	1009	4	140	2008	Delivered	2025-01-21 18:30:03.428463	2025-01-21 18:30:03.428463	

10000 rows × 10 columns

[12]: orders.dtypes

```
[12]: OrderID          int64
CustomerID        int64
OrderDateTime     datetime64[ns]
ProductID         int64
Quantity          int64
TotalAmount       int64
StoreID          int64
OrderStatus       object
ExpectedDeliveryTime  datetime64[ns]
ActualDeliveryTime  datetime64[ns]
dtype: object
```

Importing dark warehouses table

[13]: dw = pd.read_csv('Dark_Warehouses.csv')

	StoreID	StoreName	Location	Pincode	StoreCapacity
0	2000	Zepto Hub 1	Andheri West, Mumbai	400053	1134
1	2001	Zepto Hub 2	Borivali, Mumbai	400092	1225
2	2002	Zepto Hub 3	Powai, Mumbai	400076	979
3	2003	Zepto Hub 4	Koramangala, Bangalore	560095	1217
4	2004	Zepto Hub 5	Indiranagar, Bangalore	560038	1126
...
95	2095	Zepto Hub 96	Hinjewadi, Pune	411057	1028
96	2096	Zepto Hub 97	Kothrud, Pune	411038	936
97	2097	Zepto Hub 98	Viman Nagar, Pune	411014	820
98	2098	Zepto Hub 99	Vaishali Nagar, Jaipur	302021	1378
99	2099	Zepto Hub 100	Malviya Nagar, Jaipur	302017	1426

100 rows × 5 columns

[14]: dw.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   StoreID     100 non-null    int64  
 1   StoreName   100 non-null    object  
 2   Location    100 non-null    object  
 3   Pincode    100 non-null    int64  
 4   StoreCapacity 100 non-null  int64  
dtypes: int64(3), object(2)
memory usage: 4.0+ KB
```

```
[15]: dw.duplicated()
```

```
[15]: 0    False
1    False
2    False
3    False
4    False
...
95   False
96   False
97   False
98   False
99   False
Length: 100, dtype: bool
```

```
[16]: dw.dtypes
```

```
[16]: StoreID      int64
StoreName     object
Location      object
Pincode       int64
StoreCapacity int64
dtype: object
```

```
[17]: dw.isnull().sum()
```

```
[17]: StoreID      0
StoreName     0
Location      0
Pincode       0
StoreCapacity 0
dtype: int64
```

```
[18]: dw.describe()
```

```
[18]:      StoreID      Pincode  StoreCapacity
count  100.000000  100.000000  100.000000
mean  2049.500000 505902.750000 1177.290000
std   29.011492 121509.042249 203.333795
min   2000.000000 302017.000000 820.000000
25%   2024.750000 408283.500000 997.500000
50%   2049.500000 500057.500000 1211.000000
75%   2074.250000 600022.750000 1370.500000
max   2099.000000 700135.000000 1493.000000
```



Importing delivery table

```
[19]: delivery = pd.read_csv('Delivery.csv')
delivery
```

```
[19]:      DeliveryID  OrderID  DeliveryPartnerName  PickupTime  DeliveryEndTime  DistanceTraveled
0        7001      3001        Arjun Nair  04:03:4     27:03:4          1.98
1        7002      3002        Arjun Nair  00:03:4     10:03:4          7.93
2        7003      3003      Rohit Sharma  00:03:4     10:03:4          4.67
3        7004      3004      Sneha Patil  02:03:4     25:03:4          6.93
4        7005      3005      Karan Mehta  45:03:4     08:03:4          9.33
...
9995    16996     12996      Simran Kaur  46:03:4     59:03:4          8.07
9996    16997     12997      Rohit Sharma  36:03:4     54:03:4          3.61
9997    16998     12998      Karan Mehta  45:03:4     55:03:4          5.83
9998    16999     12999      Amit Singh  00:03:4     08:03:4          7.94
9999    17000     13000      Karan Mehta  22:03:4     30:03:4          8.69
```



10000 rows × 6 columns

```
[20]: delivery.dtypes
```

```
[20]: DeliveryID      int64
```

```
OrderID          int64
DeliveryPartnerName    object
PickupTime        object
DeliveryEndTime    object
DistanceTraveled float64
dtype: object
```

```
[21]: delivery.describe()
```

```
[21]:   DeliveryID  OrderID  DistanceTraveled
  count    10000.00000  10000.00000  10000.000000
  mean    12000.50000  8000.50000  5.762438
  std     2886.89568  2886.89568  2.457935
  min     7001.00000  3001.00000  1.500000
  25%    9500.75000  5500.75000  3.610000
  50%   12000.50000  8000.50000  5.740000
  75%  14500.25000  10500.25000  7.900000
  max   17000.00000  13000.00000  10.000000
```

```
[22]: delivery.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   DeliveryID      10000 non-null   int64  
 1   OrderID         10000 non-null   int64  
 2   DeliveryPartnerName 10000 non-null   object  
 3   PickupTime       10000 non-null   object  
 4   DeliveryEndTime   10000 non-null   object  
 5   DistanceTraveled 10000 non-null   float64 
dtypes: float64(1), int64(2), object(3)
memory usage: 468.9+ KB
```

```
[23]: delivery.isnull().sum()
```

```
[23]:  DeliveryID      0
  OrderID        0
  DeliveryPartnerName 0
  PickupTime      0
  DeliveryEndTime 0
  DistanceTraveled 0
  dtype: int64
```

```
[24]: delivery["PickupTime"] = pd.to_datetime(delivery["PickupTime"], format="%M:%S.%f", errors="coerce")
delivery
```

```
[24]:   DeliveryID  OrderID  DeliveryPartnerName      PickupTime  DeliveryEndTime  DistanceTraveled
  0      7001     3001      Arjun Nair  1900-01-01 00:04:03.400  27:03.4        1.98
  1      7002     3002      Arjun Nair  1900-01-01 00:00:03.400  10:03.4        7.93
  2      7003     3003     Rohit Sharma  1900-01-01 00:00:03.400  10:03.4        4.67
  3      7004     3004     Sneha Patil  1900-01-01 00:02:03.400  25:03.4        6.93
  4      7005     3005     Karan Mehta  1900-01-01 00:45:03.400  08:03.4        9.33
  ...
  9995    16996    12996     Simran Kaur  1900-01-01 00:46:03.400  59:03.4        8.07
  9996    16997    12997     Rohit Sharma  1900-01-01 00:36:03.400  54:03.4        3.61
  9997    16998    12998     Karan Mehta  1900-01-01 00:45:03.400  55:03.4        5.83
  9998    16999    12999     Amit Singh  1900-01-01 00:00:03.400  08:03.4        7.94
  9999    17000    13000     Karan Mehta  1900-01-01 00:22:03.400  30:03.4        8.69
```

10000 rows × 6 columns

```
[25]: delivery["DeliveryEndTime"] = pd.to_datetime(delivery["DeliveryEndTime"], format="%M:%S.%f", errors="coerce")
delivery #date is used only as placeholder because pandas dont have time datatype
```

```
[25]:   DeliveryID  OrderID  DeliveryPartnerName      PickupTime  DeliveryEndTime  DistanceTraveled
  0      7001     3001      Arjun Nair  1900-01-01 00:04:03.400  1900-01-01 00:27:03.400  1.98
  ...
```

1	7002	3002	Arjun Nair	1900-01-01 00:00:03.400	1900-01-01 00:10:03.400	7.93
2	7003	3003	Rohit Sharma	1900-01-01 00:00:03.400	1900-01-01 00:10:03.400	4.67
3	7004	3004	Sneha Patil	1900-01-01 00:02:03.400	1900-01-01 00:25:03.400	6.93
4	7005	3005	Karan Mehta	1900-01-01 00:45:03.400	1900-01-01 00:08:03.400	9.33
...
9995	16996	12996	Simran Kaur	1900-01-01 00:46:03.400	1900-01-01 00:59:03.400	8.07
9996	16997	12997	Rohit Sharma	1900-01-01 00:36:03.400	1900-01-01 00:54:03.400	3.61
9997	16998	12998	Karan Mehta	1900-01-01 00:45:03.400	1900-01-01 00:55:03.400	5.83
9998	16999	12999	Amit Singh	1900-01-01 00:00:03.400	1900-01-01 00:08:03.400	7.94
9999	17000	13000	Karan Mehta	1900-01-01 00:22:03.400	1900-01-01 00:30:03.400	8.69

10000 rows × 6 columns



[26]: `delivery.dtypes`

```
[26]: DeliveryID      int64
OrderID        int64
DeliveryPartnerName    object
PickupTime     datetime64[ns]
DeliveryEndTime   datetime64[ns]
DistanceTraveled float64
dtype: object
```

Importing inventory table

[27]: `inventory = pd.read_csv("Inventory.csv")`
`inventory`

	InventoryID	StoreID	ProductID	StockQuantity	ReorderLevel	LastUpdated
0	20001000	2000	1000	379	100	2025-02-19 08:46:07.994229
1	20011000	2001	1000	886	100	2025-02-19 08:46:07.994299
2	20021000	2002	1000	276	100	2025-02-19 08:46:07.994351
3	20031000	2003	1000	675	100	2025-02-18 08:46:07.994402
4	20041000	2004	1000	156	100	2025-02-20 08:46:07.994455
...
5095	20951050	2095	1050	570	100	2025-02-18 08:46:08.307134
5096	20961050	2096	1050	606	100	2025-02-17 08:46:08.307193
5097	20971050	2097	1050	685	100	2025-02-17 08:46:08.307262
5098	20981050	2098	1050	96	100	2025-02-17 08:46:08.307338
5099	20991050	2099	1050	362	100	2025-02-18 08:46:08.307423

5100 rows × 6 columns



[28]: `inventory.dtypes`

```
[28]: InventoryID      int64
StoreID        int64
ProductID      int64
StockQuantity   int64
ReorderLevel    int64
LastUpdated     object
dtype: object
```

[29]: `inventory["LastUpdated"] = pd.to_datetime(inventory["LastUpdated"], format="%Y-%m-%d %H:%M:%S.%f", errors="coerce")`
`inventory`

	InventoryID	StoreID	ProductID	StockQuantity	ReorderLevel	LastUpdated
0	20001000	2000	1000	379	100	2025-02-19 08:46:07.994229
1	20011000	2001	1000	886	100	2025-02-19 08:46:07.994299
2	20021000	2002	1000	276	100	2025-02-19 08:46:07.994351
3	20031000	2003	1000	675	100	2025-02-18 08:46:07.994402
4	20041000	2004	1000	156	100	2025-02-20 08:46:07.994455



5095	20951050	2095	1050	570	100	2025-02-18 08:46:08.307134
5096	20961050	2096	1050	606	100	2025-02-17 08:46:08.307193
5097	20971050	2097	1050	685	100	2025-02-17 08:46:08.307262
5098	20981050	2098	1050	96	100	2025-02-17 08:46:08.307338
5099	20991050	2099	1050	362	100	2025-02-18 08:46:08.307423

5100 rows × 6 columns

[30]: `inventory.dtypes`

```
[30]: InventoryID      int64
StoreID        int64
ProductID      int64
StockQuantity   int64
ReorderLevel    int64
LastUpdated     datetime64[ns]
dtype: object
```



[31]: `inventory.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5100 entries, 0 to 5099
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   InventoryID  5100 non-null   int64  
 1   StoreID      5100 non-null   int64  
 2   ProductID    5100 non-null   int64  
 3   StockQuantity 5100 non-null   int64  
 4   ReorderLevel  5100 non-null   int64  
 5   LastUpdated   5100 non-null   datetime64[ns] 
dtypes: datetime64[ns](1), int64(5)
memory usage: 239.2 KB
```

[32]: `inventory.isnull().sum()`

```
[32]: InventoryID      0
StoreID        0
ProductID      0
StockQuantity   0
ReorderLevel    0
LastUpdated     0
dtype: int64
```



[33]: `inventory.describe()`

	InventoryID	StoreID	ProductID	StockQuantity	ReorderLevel	LastUpdated
count	5.100000e+03	5100.0000	5100.000000	5100.000000	5100.0	5100
mean	2.049602e+07	2049.5000	1025.000000	526.213137	100.0	2025-02-18 20:20:09.554430720
min	2.000100e+07	2000.0000	1000.000000	50.000000	100.0	2025-02-17 08:46:07.994506
25%	2.024851e+07	2024.7500	1012.000000	293.000000	100.0	2025-02-18 08:46:07.996550400
50%	2.049602e+07	2049.5000	1025.000000	529.000000	100.0	2025-02-18 08:46:08.291908608
75%	2.074354e+07	2074.2500	1038.000000	763.000000	100.0	2025-02-19 08:46:08.299325440
max	2.099105e+07	2099.0000	1050.000000	1000.000000	100.0	2025-02-20 08:46:08.306989
std	2.886890e+05	28.8689	14.721045	272.976504	0.0	Nan

Importing products table

[34]: `products = pd.read_csv('Products.csv')`
`products`

	ProductID	ProductName	Category	Brand	PricePerUnit	StockQuantity
0	1000	Amul Taaza Milk 1L	Dairy	Amul	72	564
1	1001	Amul Butter 500g	Dairy	Amul	260	409
2	1002	Mother Dairy Curd 400g	Dairy	Mother Dairy	40	485
3	1003	Aashirvaad Atta 5kg	Grocery	Aashirvaad	255	236
4	1004	Tata Salt 1kg	Grocery	Tata	28	539
5	1005	India Gate Refined Oil 1L	Grocery	India Gate	100	226



S.	ID	Product Name	Category	Brand	Price	Stock
6	1006	Maggi Noodles 70g	Snacks	Maggi	14	936
7	1007	Lays Classic Salted 52g	Snacks	Lays	20	951
8	1008	Britannia Good Day Cashew 600g	Snacks	Britannia	150	262
9	1009	Farm Fresh Tomatoes 500g	Vegetables	Farm Fresh	35	622
10	1010	Golden Apples 1kg	Fruits	Farm Fresh	180	352
11	1011	Bananas (12 pcs)	Fruits	Farm Fresh	60	886
12	1012	Saffola Gold Oil 1L	Grocery	Saffola	180	307
13	1013	Kellogg's Corn Flakes 500g	Grocery	Kellogg's	210	519
14	1014	Dettol Handwash 750ml	Personal Care	Dettol	145	990
15	1015	Colgate Toothpaste 200g	Personal Care	Colgate	115	286
16	1016	Nescafe Classic Coffee 100g	Beverages	Nescafe	150	149
17	1017	Tata Tea Premium 500g	Beverages	Tata	240	999
18	1018	Surf Excel Matic 1kg	Grocery	Surf Excel	190	367
19	1019	Dabur Honey 500g	Grocery	Dabur	199	198
20	1020	Cadbury Dairy Milk Silk 150g	Snacks	Cadbury	160	860
21	1021	Patanjali Ghee 1L	Grocery	Patanjali	599	583
22	1022	Fortune Sunflower Oil 1L	Grocery	Fortune	160	402
23	1023	Haldiram's Bhujia 400g	Snacks	Haldiram's	135	275
24	1024	Bournvita 500g	Beverages	Cadbury	225	719
25	1025	Dove Soap 125g	Personal Care	Dove	45	759
26	1026	Gillette Mach3 Razor	Personal Care	Gillette	299	847
27	1027	Real Mixed Fruit Juice 1L	Beverages	Real	99	286
28	1028	Catch Black Pepper 100g	Grocery	Catch	150	325
29	1029	Horlicks Classic Malt 1kg	Beverages	Horlicks	385	879
30	1030	Johnson's Baby Powder 400g	Personal Care	Johnson's	175	953
31	1031	Himalaya Face Wash 100ml	Personal Care	Himalaya	130	822
32	1032	Whisper Ultra Pads (30 pcs)	Personal Care	Whisper	320	561
33	1033	Oreo Chocolate Biscuits 300g	Snacks	Oreo	85	898
34	1034	Lizol Floor Cleaner 1L	Grocery	Lizol	210	427
35	1035	Sprite 2L	Beverages	Sprite	90	925
36	1036	Frooti Mango Drink 1L	Beverages	Frooti	65	946
37	1037	Dabur Chyawanprash 1kg	Health	Dabur	350	612
38	1038	Coca-Cola 2L	Beverages	Coca-Cola	90	789
39	1039	Eveready AA Batteries (4 pcs)	Electronics	Eveready	110	100
40	1040	Sensodyne Toothpaste 100g	Personal Care	Sensodyne	160	566
41	1041	Nestle Everyday Milk Powder 500g	Dairy	Nestle	275	977
42	1042	Kissan Mixed Fruit Jam 500g	Grocery	Kissan	180	819
43	1043	Good Day Butter Cookies 600g	Snacks	Britannia	140	113
44	1044	Brooke Bond Red Label 500g	Beverages	Brooke Bond	170	447
45	1045	Parle-G Biscuits 800g	Snacks	Parle	80	146
46	1046	Sprite 2L	Beverages	Sprite	90	338
47	1047	Head & Shoulders Shampoo 650ml	Personal Care	Head & Shoulders	385	635
48	1048	Real Orange Juice 1L	Beverages	Real	110	159
49	1049	Gillette Shaving Foam 200ml	Personal Care	Gillette	240	237
50	1050	Dabur Red Toothpaste 200g	Personal Care	Dabur	99	446

[35]: products.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 51 entries, 0 to 50
```

```
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype  
 --- 
 0   ProductID   51 non-null     int64  
 1   ProductName  51 non-null     object  
 2   Category     51 non-null     object  
 3   Brand        51 non-null     object  
 4   PricePerUnit 51 non-null     int64  
 5   StockQuantity 51 non-null     int64  
 dtypes: int64(3), object(3) 
 memory usage: 2.5+ KB
```

```
[36]: products.dtypes
```

```
[36]: ProductID      int64
ProductName    object
Category      object
Brand         object
PricePerUnit   int64
StockQuantity  int64
dtype: object
```

```
[37]: products.isnull().sum()
```

```
[37]: ProductID      0
ProductName    0
Category      0
Brand         0
PricePerUnit   0
StockQuantity  0
dtype: int64
```

```
[38]: products.describe()
```

	ProductID	PricePerUnit	StockQuantity
count	51.000000	51.000000	51.000000
mean	1025.000000	167.333333	551.647059
std	14.866069	109.806314	285.374899
min	1000.000000	14.000000	100.000000
25%	1012.500000	90.000000	296.500000
50%	1025.000000	150.000000	539.000000
75%	1037.500000	210.000000	834.500000
max	1050.000000	599.000000	999.000000

Importing Returns table

```
[39]: returns = pd.read_csv('Returns.csv')
returns
```

	ReturnID	OrderID	ProductID	CustomerID	StoreID	ReturnReason	ReturnDate	RefundAmount	ReturnStatus
0	17794	12794	1011	15794	2075	Stale Product	2025-02-12 08:33:03.428463	300	Processed
1	13898	8898	1030	11898	2050	Damaged Item	2025-01-25 20:43:03.428463	875	Processed
2	17271	12271	1012	15271	2056	Damaged Item	2025-02-08 22:13:03.428463	360	Processed
3	12959	7959	1028	10959	2050	Stale Product	2025-02-11 06:23:03.428463	600	Processed
4	13863	8863	1024	11863	2016	Stale Product	2025-02-22 05:21:03.428463	1125	Processed
...
995	16028	11028	1010	14028	2004	Stale Product	2025-01-28 21:42:03.428463	360	Processed
996	8750	3750	1023	6750	2003	Damaged Item	2025-02-22 01:05:03.428463	675	Processed
997	17265	12265	1035	15265	2059	Stale Product	2025-02-06 18:15:03.428463	450	Processed
998	15520	10520	1044	13520	2087	Damaged Item	2025-02-22 18:00:03.428463	850	Processed
999	16852	11852	1028	14852	2008	Damaged Item	2025-01-31 14:37:03.428463	300	Processed

1000 rows × 9 columns

```
[40]: returns.dtypes
```

```
[40]: ReturnID      int64
OrderID       int64
```

```
ProductID      int64
CustomerID     int64
StoreID        int64
ReturnReason    object
ReturnDate      object
RefundAmount    int64
ReturnStatus    object
dtype: object
```

```
[41]: returns['ReturnDate'] = pd.to_datetime(returns['ReturnDate'], format = '%Y-%m-%d %H:%M:%S.%f')
returns
```

```
[41]:
```

	ReturnID	OrderID	ProductID	CustomerID	StoreID	ReturnReason	ReturnDate	RefundAmount	ReturnStatus
0	17794	12794	1011	15794	2075	Stale Product	2025-02-12 08:33:03.428463	300	Processed
1	13898	8898	1030	11898	2050	Damaged Item	2025-01-25 20:43:03.428463	875	Processed
2	17271	12271	1012	15271	2056	Damaged Item	2025-02-08 22:13:03.428463	360	Processed
3	12959	7959	1028	10959	2050	Stale Product	2025-02-11 06:23:03.428463	600	Processed
4	13863	8863	1024	11863	2016	Stale Product	2025-02-22 05:21:03.428463	1125	Processed
...
995	16028	11028	1010	14028	2004	Stale Product	2025-01-28 21:42:03.428463	360	Processed
996	8750	3750	1023	6750	2003	Damaged Item	2025-02-22 01:05:03.428463	675	Processed
997	17265	12265	1035	15265	2059	Stale Product	2025-02-06 18:15:03.428463	450	Processed
998	15520	10520	1044	13520	2087	Damaged Item	2025-02-22 18:00:03.428463	850	Processed
999	16852	11852	1028	14852	2008	Damaged Item	2025-01-31 14:37:03.428463	300	Processed

1000 rows × 9 columns

```
[42]: returns.dtypes
```

```
[42]:
```

	ReturnID	int64
OrderID		int64
ProductID		int64
CustomerID		int64
StoreID		int64
ReturnReason		object
ReturnDate		datetime64[ns]
RefundAmount		int64
ReturnStatus		object
dtype: object		

```
[43]: returns.isnull().sum()
```

```
[43]:
```

	ReturnID	0
OrderID		0
ProductID		0
CustomerID		0
StoreID		0
ReturnReason		0
ReturnDate		0
RefundAmount		0
ReturnStatus		0
dtype: int64		

```
[44]: returns.describe()
```

```
[44]:
```

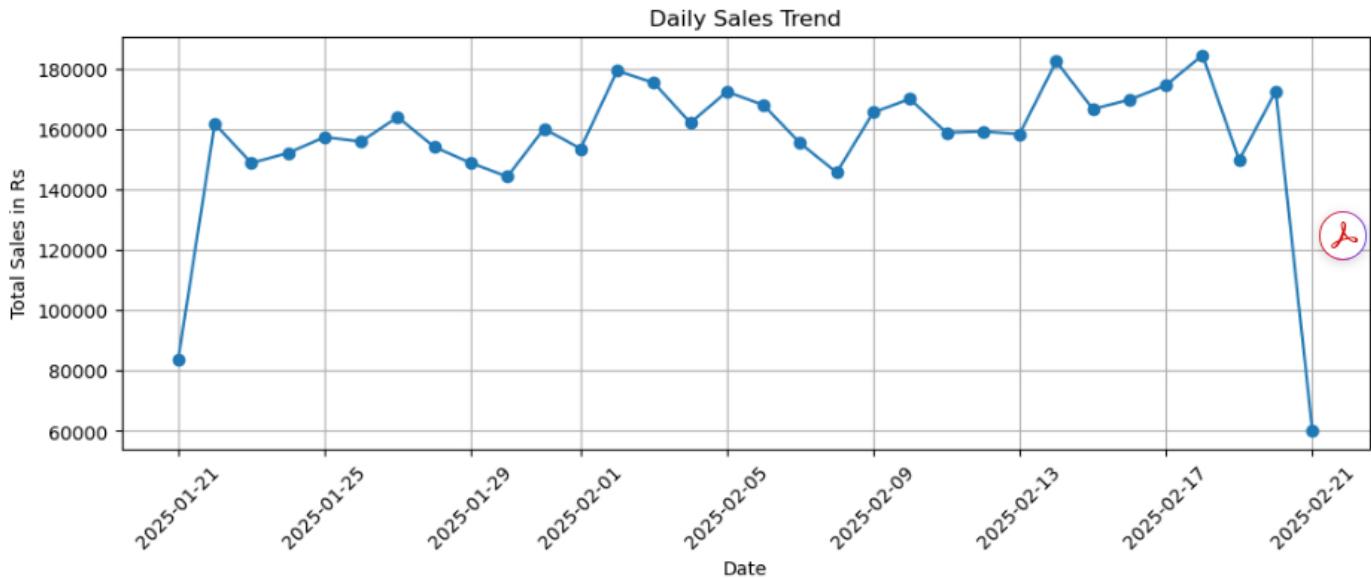
	ReturnID	OrderID	ProductID	CustomerID	StoreID	ReturnDate	RefundAmount	
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000		1000	1000.000000
mean	13109.208000	8109.208000	1024.361000	11109.208000	2047.924000	2025-02-06 22:12:09.668462848	479.203000	
min	8008.000000	3008.000000	1000.000000	6008.000000	2000.000000	2025-01-22 02:08:03.428463	14.000000	
25%	10633.500000	5633.500000	1011.000000	8633.500000	2021.000000	2025-01-29 23:54:03.428463104	180.000000	
50%	13029.500000	8029.500000	1024.000000	11029.500000	2048.000000	2025-02-07 03:07:03.428463104	360.000000	
75%	15723.750000	10723.750000	1037.000000	13723.750000	2073.000000	2025-02-13 22:10:48.428463104	675.000000	
max	17987.000000	12987.000000	1050.000000	15987.000000	2099.000000	2025-02-23 03:04:03.428463	2995.000000	
std	2901.123785	2901.123785	14.731113	2901.123785	29.178273		Nan	419.752689

What is the sales trend over time (day/week/month)?

```
[45]: # Group by date (ignoring time part)
daily_sales = orders.groupby(orders['OrderDateTime'].dt.date)['TotalAmount'].sum().reset_index()

# Rename columns for clarity
daily_sales.columns = ['Date', 'TotalSales']

# Plot line chart
plt.figure(figsize = (12,4))
plt.plot(daily_sales['Date'], daily_sales['TotalSales'], marker = 'o')
plt.title('Daily Sales Trend')
plt.xlabel('Date')
plt.ylabel('Total Sales in Rs')
plt.xticks(rotation = 45)
plt.grid(True)
plt.show()
```



```
[46]: # Group sales by day of week
daily_sales = orders.groupby(orders['OrderDateTime'].dt.day_name())['TotalAmount'].sum().reset_index()

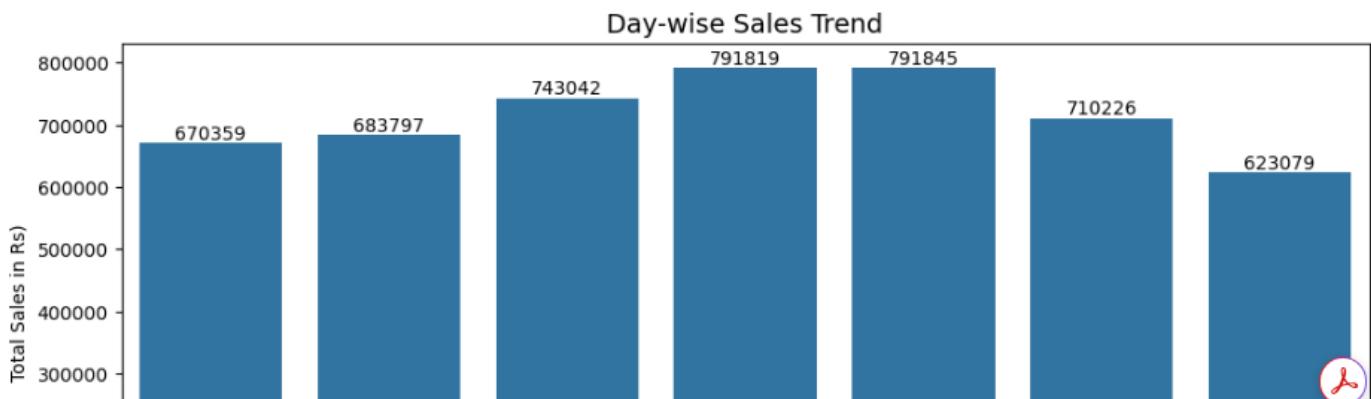
# Rename columns
daily_sales.columns = ['Day', 'TotalSales']

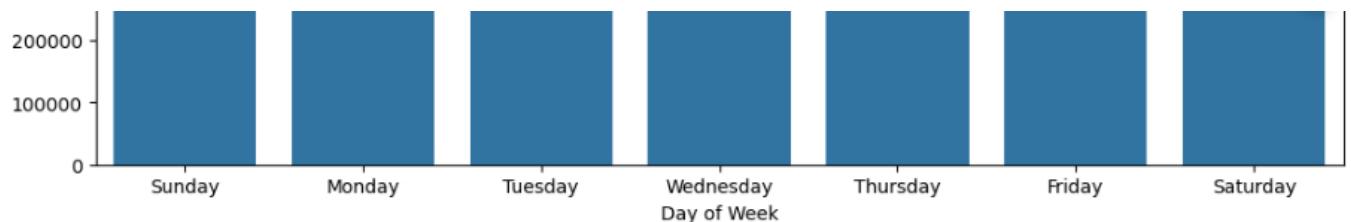
# Define correct order for days
day_order = ['Sunday', 'Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday']
daily_sales['Day'] = pd.Categorical(daily_sales['Day'], categories=day_order, ordered=True)
daily_sales = daily_sales.sort_values('Day')

# Plot using seaborn
plt.figure(figsize=(12,5))
ax = sns.barplot(x='Day', y='TotalSales', data=daily_sales)

# Add Labels on bars
for container in ax.containers:
    ax.bar_label(container, fmt='%d', label_type='edge')

# Chart formatting
plt.title('Day-wise Sales Trend', fontsize=14)
plt.xlabel('Day of Week')
plt.ylabel('Total Sales in Rs')
plt.show()
```





Which product categories generate the most revenue?

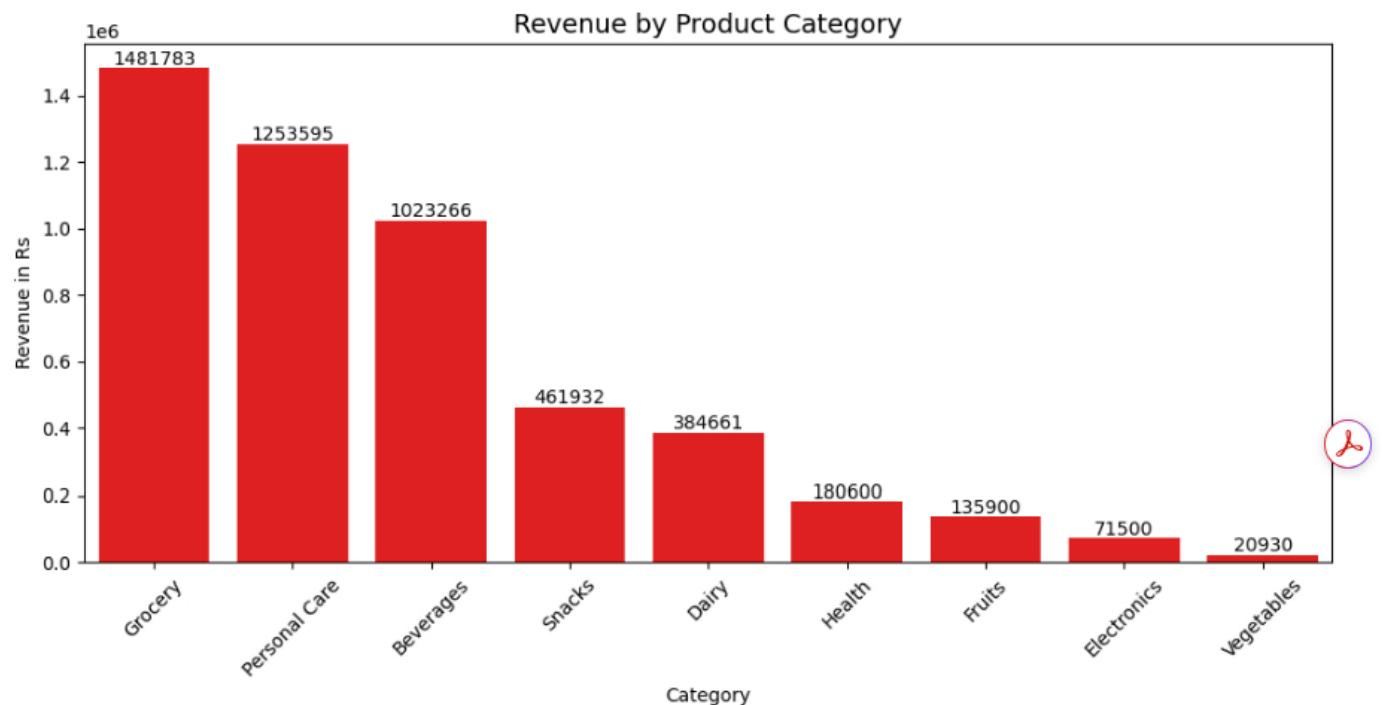
```
[47]: # Group by category and calculate revenue
category_sales = orders.merge(products, on="ProductID") \
    .groupby("Category")["TotalAmount"].sum() \
    .reset_index()

# Sort by revenue for better visuals
category_sales = category_sales.sort_values("TotalAmount", ascending=False)

# Plot
plt.figure(figsize=(12,5))
ax = sns.barplot(x="Category", y="TotalAmount", data=category_sales, color="red")

# Add labels on bars
for container in ax.containers:
    ax.bar_label(container, fmt="%d", fontsize=10)

# Formatting
plt.title("Revenue by Product Category", fontsize=14)
plt.xlabel("Category")
plt.ylabel("Revenue in Rs")
plt.xticks(rotation=45)
plt.show()
```



Which brands/products are the top sellers?

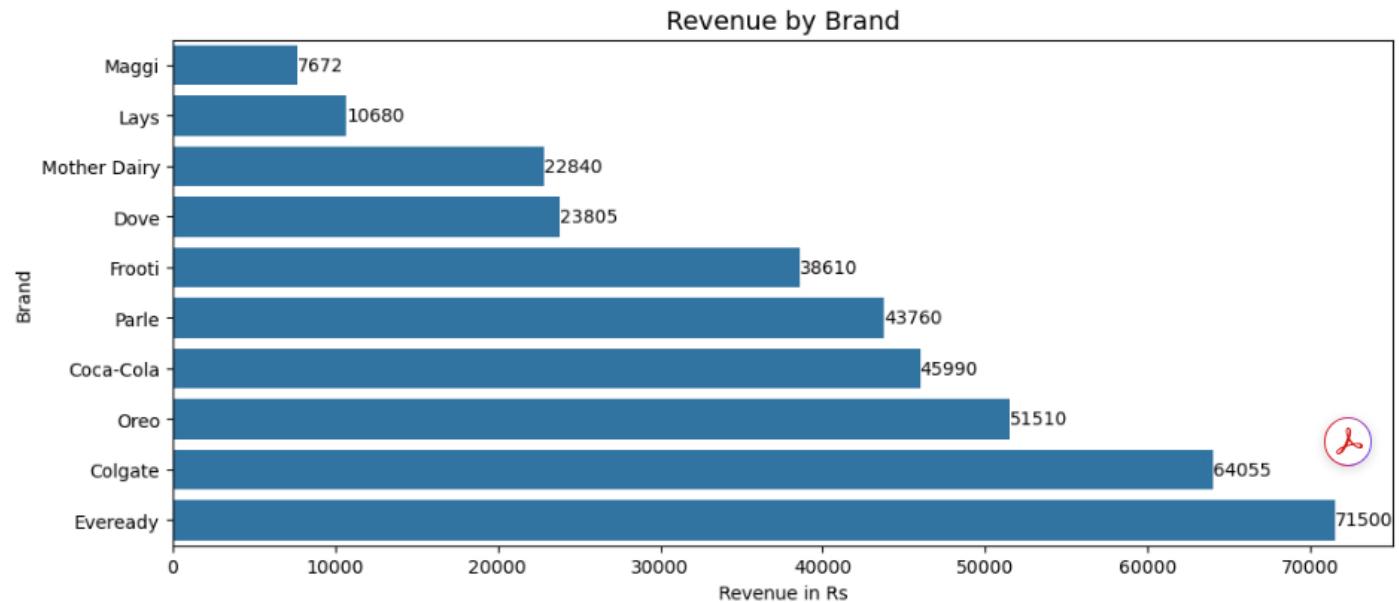
```
[48]: # Group by category and calculate revenue
brand_sales = products.merge(orders, on='ProductID')\ 
    .groupby('Brand')['TotalAmount'].sum()\ 
    .reset_index()

# Sort by revenue for better visuals
brand_sales = brand_sales.sort_values('TotalAmount').head(10)

# Plot
plt.figure(figsize=(12,5))
ax = sns.barplot(x='TotalAmount', y='Brand', data=brand_sales, orient='h')
```

```
# Add Labels on bars
for container in ax.containers:
    ax.bar_label(container, fmt="%d")

# Formatting
plt.title("Revenue by Brand", fontsize=14)
plt.xlabel("Revenue in Rs")
plt.ylabel("Brand")
plt.show()
```



Revenue contribution by each store/dark warehouse

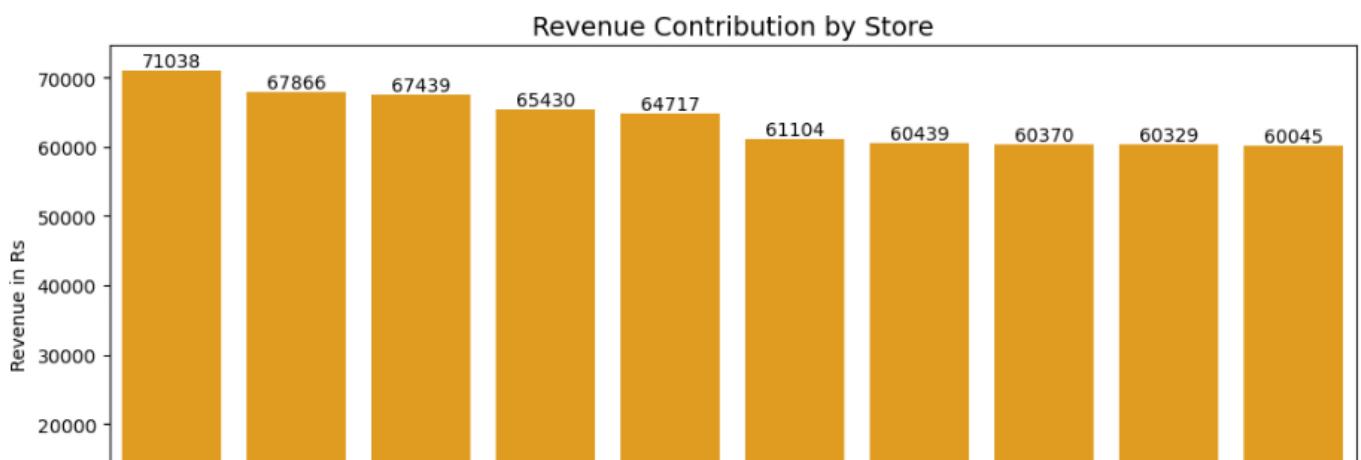
```
[49]: # Group by category and calculate revenue
store_sales = dw.merge(orders, on='StoreID')\
    .groupby('StoreName')['TotalAmount'].sum() \
    .reset_index()

# Sort for better visuals
store_sales = store_sales.sort_values('TotalAmount', ascending = False).head(10)

#plot
plt.figure(figsize=(12,5))
ax = sns.barplot(x='StoreName', y='TotalAmount', data = store_sales, color ='orange')

# Add Labels on bars
for container in ax.containers:
    ax.bar_label(container, fmt="%d")

# Formatting
plt.title("Revenue Contribution by Store", fontsize=14)
plt.xlabel("Store Name")
plt.ylabel("Revenue in Rs")
plt.xticks(rotation=45)
plt.show()
```





Which store has the highest sales per square capacity?

```
[50]: # Group by category and calculate revenue
store_capacity = dw.merge(orders, on='StoreID')\
    .groupby(['StoreName', 'StoreCapacity'])['TotalAmount'].sum()\
    .reset_index()

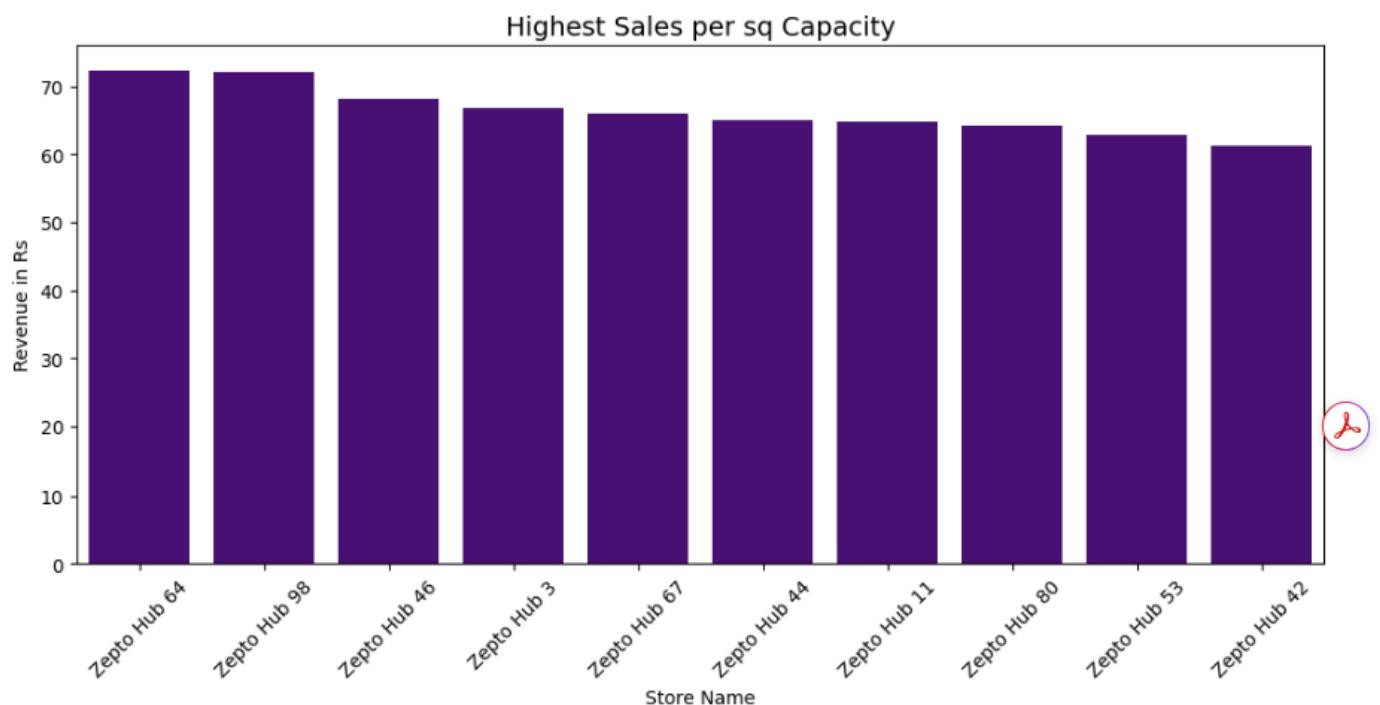
# Calculate revenue per capacity
store_capacity["RevenuePerCapacity"] = store_capacity["TotalAmount"] / store_capacity["StoreCapacity"]

# Sort by ratio and take top 10
store_capacity = store_capacity.sort_values("RevenuePerCapacity", ascending=False).head(10)

#plot
plt.figure(figsize =(12,5))
ax = sns.barplot(x = 'StoreName', y= 'RevenuePerCapacity', data = store_capacity, color = 'indigo')

# Add labels on bars
for container in ax.containers:
    ax.bar_label(container, fmt="%2f")

# Formatting
plt.title("Highest Sales per sq Capacity", fontsize=14)
plt.xlabel("Store Name")
plt.ylabel("Revenue in Rs")
plt.xticks(rotation=45)
plt.show()
```



Total stock quantity available per store

```
[51]: # Group by category and calculate revenue
store = inventory.merge(dw, on='StoreID')\
    .groupby('StoreName')['StockQuantity'].sum().reset_index()

#sort
store = store.sort_values('StockQuantity').head(20)
```

```
#plot
plt.figure(figsize =(12,5))
sns.lineplot(data=store, x="StoreName", y="StockQuantity", marker ='o', color = 'green')

# Formatting
plt.title("Total Quantity Stocks Available", fontsize=14)
plt.xlabel("Store Name")
plt.ylabel("Revenue in Rs")
plt.xticks(rotation=90)
plt.grid()
plt.show()
```



On-time vs Late Deliveries (overall)

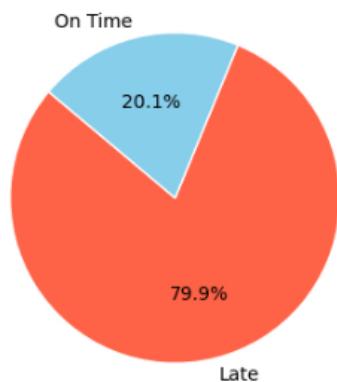
```
[71]: # Classify deliveries as On-time or Late
orders["DeliveryStatus"] = orders.apply(
    lambda x: "Late" if x["ActualDeliveryTime"] > x["ExpectedDeliveryTime"] else "On Time",
    axis=1
)

# Count status
delivery_status = orders["DeliveryStatus"].value_counts()

# Plot pie chart
plt.figure(figsize=(4,4))
plt.pie(delivery_status, labels=delivery_status.index, autopct="%1.1f%%",
        colors=["tomato", "skyblue"], startangle=140, wedgeprops={'edgecolor':'white'})

plt.title("On-Time vs Late Deliveries", fontsize=14)
plt.show()
```

On-Time vs Late Deliveries



Delivery distance distribution

```
[69]: # Plot histogram
plt.figure(figsize=(12,5))
sns.histplot(delivery["DistanceTraveled"], bins=40, kde=True, color="skyblue")

# Formatting
plt.title("Delivery Distance Distribution", fontsize=14)
plt.xlabel("Distance Traveled in KM")
plt.ylabel("Number of Deliveries")
plt.grid(axis="y", linestyle="--", alpha=0.7)
plt.show()
```



Top 5 most returned products.

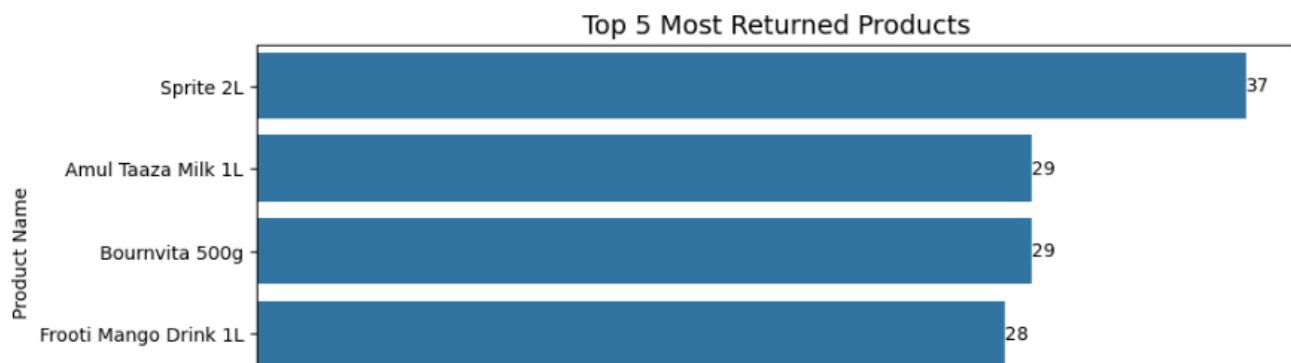
```
[82]: # Merge products with returns
product_returns = returns.merge(products, on="ProductID")\
    .groupby('ProductName')['ReturnID'].count().reset_index()

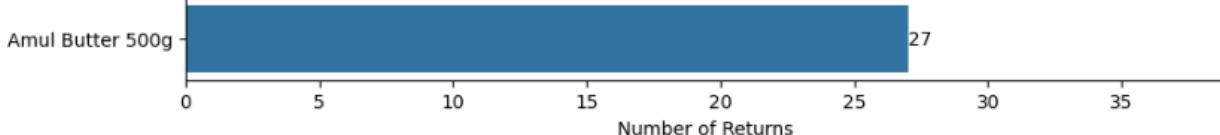
# Sort and take top 5
top_returns = product_returns.sort_values("ReturnID", ascending=False).head(5)

# Plot
plt.figure(figsize=(10,4))
ax = sns.barplot(x="ReturnID", y="ProductName", data=top_returns)

# Add labels on bars
for container in ax.containers:
    ax.bar_label(container, fmt='%d')

# Formatting
plt.title("Top 5 Most Returned Products", fontsize=14)
plt.xlabel("Number of Returns")
plt.ylabel("Product Name")
plt.show()
```





Return rate per category

```
[104]: # Merge returns with products & orders with products to get category
returns_products = returns.merge(products, on = 'ProductID')

orders_products = orders.merge(products, on = 'ProductID')

# Total returns & orders per category
returns_per_category = returns_products.groupby('Category')['ReturnID'].count().reset_index(name="TotalReturns")
orders_per_category = orders_products.groupby('Category')['OrderID'].count().reset_index(name="TotalOrders")

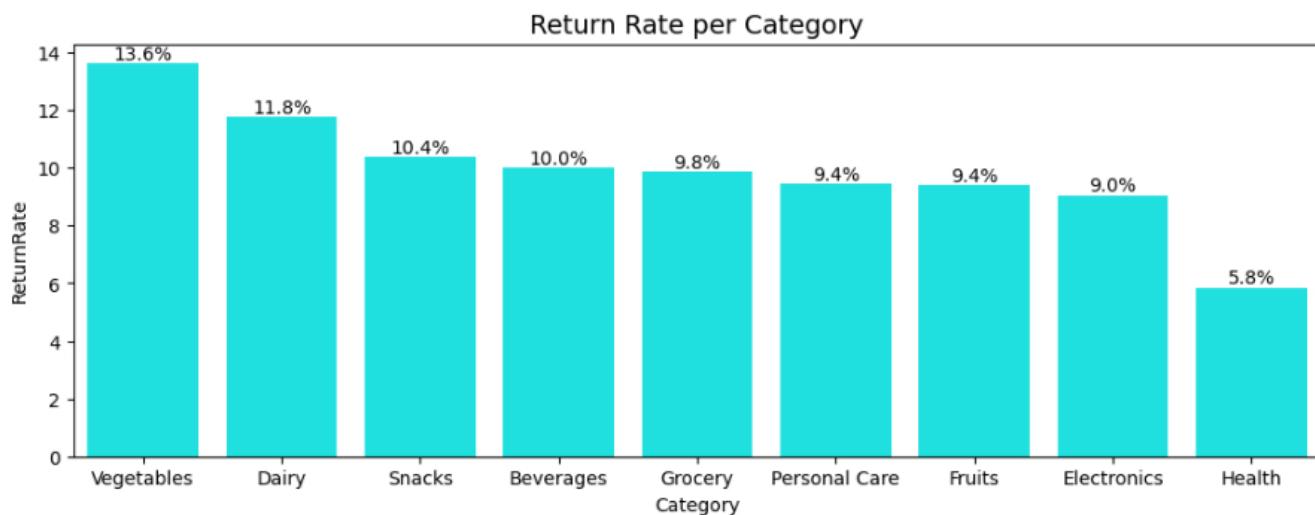
# Merge and calculate return rate %
return_rate = returns_per_category.merge(orders_per_category, on="Category")
return_rate["ReturnRate"] = (return_rate["TotalReturns"] / return_rate["TotalOrders"]) * 100

#sort
return_rate = return_rate.sort_values("ReturnRate", ascending = False)

#plot
plt.figure(figsize= (12,4))
ax = sns.barplot(data = return_rate, x ='Category', y='ReturnRate', color='cyan')

#labels
for container in ax.containers:
    ax.bar_label(container, fmt='%.1f%%')

#Formatting
plt.title("Return Rate per Category", fontsize=14)
plt.show()
```



[1]:

⟳ ⌂ ⌄ ⌅ ⌆ ⌇ ⌈ ⌉ ⌊ ⌋

